

The fault in our traits & the borrow checker

Aman Sharma

Aftershoot Inc.

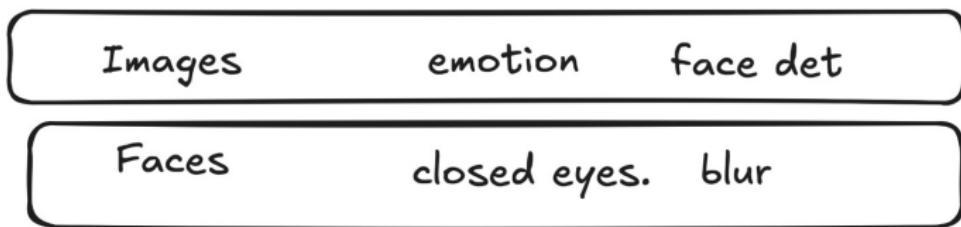
January 15, 2026

Now I was trying to create an encapsulation of domain model into traits.

Culling is selecting of ?good? images. The definition of good/bad is judged by multiple factors, but for the sake of this presentation we take every image, then run ML models on them get scores, and run algorithm on them to get selection/rejections.

a small & wrong diagram of our code structure looks like this,
remind you this is of only one pipeline, each pipeline has a separate
structure like this, this causes massive code duplication.

Wedding Engagement



This in term made very brittle and hard to reason about, any change and addition to the feature or code required writing something horrendous like this, or simply use a dozen of macros and work in together.

```
#[pilot]
impl<ICnstr, F> Pilot for FaceDupLayerOptIn<ICnstr, F>
where
    ICnstr: PartialImageConstructor,
    F: Fn(&PartialImage<ICnstr>) → ndarray::Array1<f32>,
    <ICnstr as PartialFaceConstructor>::PartialFaceScores:
        FaceScoresFaceDupFuxx + FaceScoresFuxx + FaceScoresImportanceChangeFuxx,
    <ICnstr as PartialImageConstructor>::PartialImageProperties: ImagePropsFuxxFaceDup,
    <ICnstr as PartialImageConstructor>::PartialImageScores:
        ImageScoresFuxxFaceDup + ImageScoresFuxxPerson,
{
    type Input<'i> = Vec<Vec<PartialImage<ICnstr>>>;
    type Output<'o> = Vec<Vec<PartialImage<ICnstr>>>;
}
```

Even with the macros, composing features was difficult it is *also* because of the how quickly changes were made in the pipeline, and also basically the issue being of data drilling, as in data that is created on end of a program is required on the other end of program, also testing was also impossible.

The core reason of creating any type of encapsulation is to have less code duplication, traits are good for code generation.

But they force you to think in term of little small pieces of data/methods, that can be accessed from within the traits this severly restricts with what a computer can do, it primarily restricts things that look at two ?random? pieces of data, that might not have any explicit relationship in the domain model but an implicit relationship at runtime, this puts a bizzare restriction on a given software. that is how if you want to do things within an pipeline, everything works fine-ish, but accross even the same logic need to respresented again and again.

What's the solution??

We can flip the encapsulation boundary around the feature(s) and not the domain model itself.

instead of having traits that define methods and GATs, we directly do the thing that we want to do.

```
1 pub struct CullingGlobalState {  
2     pub scores: ImageScores,  
3     pub properties: ImageProperties,  
4     pub images: Vec<Image>,  
5 }  
6  
6 pub struct ImageScores {  
7     pub blur: Vec<f32>  
8 }  
9  
10 pub struct ImageProperties {  
11     pub is_selected: Vec<bool>  
12 }  
13
```

this kind of similar to how a ECS would work, this makes the code easy to reason about, and no encapsulation would never bottleneck the application development.

So where does borrow checker fit in all of this?

Well, we tend to think of memory being two kind of memory, stack vs heap

Rust was created as a replacement of CPP. "Compile-time hierarchy of encapsulation that matches the domain model" - Casey Muratori